

# Comparative Analysis of Reinforcement Learning Algorithms for Autonomous Blue Team Defense in CybORG

Research Study on CybORG Blue Agent Training

Vasanth Iyer

*Based on the CybORG Framework*

*Standen et al., IJCAI-21 1st International Workshop on Adaptive Cyber Defense*

February 2026

## Abstract

This study presents a comprehensive comparison of multiple reinforcement learning algorithms for training autonomous Blue team (defender) agents in the CybORG (Cyber Operations Research Gym) environment. We evaluate Proximal Policy Optimization (PPO), Long Short-Term Memory PPO (LSTM-PPO), and Deep Q-Network (DQN) approaches on Scenario 1b with a hierarchical action space. Our experiments reveal that DQN significantly outperforms policy gradient methods in learning state-dependent defensive strategies, achieving 39 action changes per episode compared to 0 for PPO-based methods. The DQN agent learned a cyclic defense pattern utilizing Monitor, Analyse, Remove, and Restore actions across multiple hosts, demonstrating genuine reactive behavior to the Red agent's attack sequence.

## 1 Introduction

Autonomous Cyber Operations (ACO) represents a critical challenge in modern cybersecurity, requiring intelligent agents capable of defending computer networks against sophisticated adversaries. The CybORG framework [1] provides a standardized environment for developing and evaluating such agents through both simulation and emulation modes.

This work focuses on training Blue team (defender) agents against the deterministic B\_lineAgent Red team attacker in CybORG's Scenario 1b. We investigate multiple reinforcement learning approaches and their effectiveness in learning state-dependent defensive policies.

### 1.1 Problem Statement

The primary challenge identified during this research was **action collapse** - where trained agents converge to repeatedly selecting a single action regardless of the environment state. This phenomenon was observed across multiple algorithm configurations and represents a significant barrier to developing effective autonomous defenders.

### 1.2 Contributions

- Comprehensive comparison of PPO, LSTM-PPO, and DQN for Blue team training
- Identification of the action collapse problem in policy gradient methods
- Demonstration that DQN learns genuine state-dependent policies
- Hierarchical action space design reducing complexity from 54 to 52 actions
- Detailed analysis of learned defensive strategies

## 2 Background

### 2.1 CybORG Framework

CybORG [1] is a gym for Autonomous Cyber Operations research featuring:

- **Simulation Mode:** Finite state machine representation for rapid training
- **Emulation Mode:** AWS-based virtual infrastructure for validation
- **OpenAI Gym Interface:** Standard RL interaction protocol
- **Adversarial Scenarios:** Red vs Blue team competitions

The original CybORG paper demonstrated successful training of Red team agents using Deep Q-Networks with LSTM (DRQN), achieving a 66% transfer rate from simulation to emulation. Our work extends this to Blue team defense.

### 2.2 Scenario 1b Configuration

Table 1: Scenario 1b Network Configuration

Subnet	Hosts	Role
User Subnet	User0, User1, User2, User3, User4	Entry points
Enterprise Subnet	Enterprise0, Enterprise1, Enterprise2	Mid-tier targets
Operational Subnet	Op_Host0, Op_Host1, Op_Host2, Op_Server0	Critical assets
Defense	Defender	Blue team base

### 2.3 Action Space

Blue team agents have access to the following action types:

Table 2: Blue Team Action Types

Action	Description
Sleep	No operation
Monitor	Network-wide surveillance for threats
Analyse <host>	Detailed inspection of specific host
Remove <host>	Remove malicious processes/files
Restore <host>	Full system restore (expensive)
Misinform <host>	Deploy deception/honeypots

### 2.4 Red Team: B\_lineAgent

The B\_lineAgent follows a deterministic attack pattern:

1. Discover subnet via scanning
2. Exploit User subnet hosts
3. Pivot to Enterprise subnet
4. Escalate privileges toward Op\_Server0

This predictable sequence should theoretically enable Blue agents to learn anticipatory defenses.

### 3 Methodology

#### 3.1 Hierarchical Action Space

To reduce action space complexity, we implemented a hierarchical encoding:

$$a = t \times |H| + h \quad (1)$$

where  $t \in \{0, 1, 2, 3\}$  represents action type (Monitor, Analyse, Remove, Restore),  $h \in \{0, \dots, 12\}$  represents host index, and  $|H| = 13$  is the number of hosts.

This reduces the action space from 54 discrete actions to 52 ( $4 \times 13$ ), removing Sleep and Misinform.

#### 3.2 Algorithms Evaluated

##### 3.2.1 Proximal Policy Optimization (PPO)

PPO [2] is a policy gradient method that optimizes:

$$L^{CLIP}(\theta) = \mathbb{E}_t \left[ \min(r_t(\theta)\hat{A}_t, \text{clip}(r_t(\theta), 1 - \epsilon, 1 + \epsilon)\hat{A}_t) \right] \quad (2)$$

##### Configuration:

- Learning rate:  $3 \times 10^{-4}$
- Entropy coefficient: 0.05
- Network: MLP [256, 256]
- Training steps: 500,000

##### 3.2.2 LSTM-PPO (Recurrent PPO)

RecurrentPPO extends PPO with LSTM memory to capture temporal patterns:

##### Configuration:

- LSTM hidden size: 128
- LSTM layers: 1
- Shared LSTM for policy and value
- Entropy coefficient: 0.01
- Training steps: 500,000

##### 3.2.3 Deep Q-Network (DQN)

DQN [3] learns action-value function  $Q(s, a)$  directly:

$$Q(s, a) = r + \gamma \max_{a'} Q(s', a') \quad (3)$$

##### Configuration:

- Learning rate:  $1 \times 10^{-4}$
- Replay buffer: 100,000 transitions
- Exploration:  $\epsilon$  from 1.0 to 0.05 over 30% of training

- Network: MLP [256, 256, 128]
- Target update interval: 1,000 steps
- Training steps: 500,000

### 3.3 Evaluation Metrics

- **Mean Reward:** Average episode return (native CybORG reward)
- **Unique Actions:** Number of distinct actions per episode
- **Top Action %:** Percentage of most common action
- **Action Changes:** Number of times action changes within episode
- **Deterministic vs Stochastic:** Comparing arg max vs sampled action selection

## 4 Results

### 4.1 Overall Performance Comparison

Table 3: Deterministic Evaluation Results

Algorithm	Mean Reward	Unique Actions	Top Action %	State-Dependent
PPO (Flat)	-20.0	1	100.0%	No
PPO (Hierarchical)	-291.4	1	100.0%	No
LSTM-PPO	-1115.8	1	100.0%	No
<b>DQN</b>	<b>-228.1</b>	<b>6</b>	<b>56.7%</b>	<b>Yes</b>

### 4.2 Stochastic vs Deterministic Performance

Table 4: PPO Hierarchical: Stochastic vs Deterministic

Metric	Deterministic	Stochastic
Mean Reward	-291.4	-163.5
Unique Actions/Episode	1.0	26.8
Total Unique Actions	1	39
Most Common Action %	100.0%	26.1%

This reveals that PPO learns a **flat probability distribution** - stochastic sampling shows diversity, but deterministic arg max collapses to a single action.

### 4.3 Action Type Distribution

Table 5: Action Type Usage (Deterministic Mode)

Algorithm	Monitor	Analyse	Remove	Restore
PPO Flat	0%	0%	0%	100%
PPO Hierarchical	0%	0%	0%	100%
LSTM-PPO	0%	100%	0%	0%
<b>DQN</b>	<b>2.0%</b>	<b>74.8%</b>	<b>7.9%</b>	<b>15.3%</b>

DQN is the only algorithm utilizing multiple action types, with a preference for Analyse operations.

### 4.4 State-Dependency Analysis

Table 6: Action Changes per 50-Step Episode

Algorithm	Action Changes	Classification
PPO Flat	0	Static
PPO Hierarchical	0	Static
LSTM-PPO	0	Static
<b>DQN</b>	<b>39</b>	<b>State-Dependent</b>

### 4.5 DQN Learned Policy

Analysis of Q-values reveals DQN learned meaningful preferences:

Table 7: DQN Q-Values by Action Type (Initial State)

Action Type	Mean Q	Max Q	Min Q
Monitor	-2.97	-2.77	-3.59
Analyse	-2.92	-2.79	-3.30
Remove	-2.95	-2.81	-3.30
Restore	-3.91	-3.60	-4.39

Key insight: DQN learned to **avoid Restore** (lowest Q-values) and prefer Analyse/Remove/Monitor.

#### 4.5.1 Learned Defense Cycle

DQN learned a repeating defensive pattern:

1. **Step 0:** Monitor (reconnaissance)
2. **Steps 1-2:** Analyse Op\_Server0 (check critical server)
3. **Step 3:** Remove Enterprise1 (clean compromised host)
4. **Steps 4+:** Alternate between Analyse and Remove

Table 8: DQN Top 6 Actions (Deterministic Evaluation)

Action	Count	Percentage
Analyse Op_Server0	567	56.7%
Analyse Op_Host0	156	15.6%
Restore Op_Server0	153	15.3%
Remove Enterprise1	79	7.9%
Analyse User0	25	2.5%
Monitor	20	2.0%

## 5 Discussion

### 5.1 Why Policy Gradient Methods Failed

PPO and LSTM-PPO learn a policy  $\pi(a|s)$  that produces a probability distribution over actions. Our experiments show these methods converge to near-uniform distributions:

$$\pi(a|s) \approx \frac{1}{|A|} \quad \forall a \in A \quad (4)$$

When using deterministic inference ( $\arg \max$ ), tiny numerical differences determine the selected action, resulting in the same action regardless of state.

#### Root causes:

1. High entropy coefficients encourage exploration but prevent commitment
2. Observations may lack sufficient discriminative information
3. Action space too large relative to observation signal

### 5.2 Why DQN Succeeded

DQN learns  $Q(s, a)$  values directly, where each action's expected return is estimated independently. This architecture naturally produces **state-dependent** action selection:

$$a^* = \arg \max_a Q(s, a) \quad (5)$$

Key advantages:

1. Q-values differentiate actions based on state features
2. Replay buffer provides stable, decorrelated training
3.  $\epsilon$ -greedy exploration is simpler than entropy regularization
4. No distribution collapse - each Q-value is learned independently

### 5.3 Comparison with Original CybORG Paper

Table 9: Comparison with Standen et al. (2021)

Aspect	Original Paper	This Work
Agent Type	Red (attacker)	Blue (defender)
Algorithm	DRQN (DQN + LSTM)	DQN, PPO, LSTM-PPO
Scenario	3-host penetration	Scenario 1b (13 hosts)
Action Space	8 actions	52 actions (hierarchical)
Training Steps	2,500 iterations	500,000 steps
Success Metric	System access	Reward maximization
Transfer Rate	66% (sim → emu)	N/A (simulation only)

The original CybORG paper demonstrated DQN with LSTM (DRQN) for Red team training. Our results confirm that value-based methods (DQN) are more effective than policy gradient methods (PPO) for this domain, extending the finding to Blue team defense.

### 5.4 Limitations

1. **Simulation Only:** No emulation validation performed
2. **Single Red Agent:** Only tested against B\_lineAgent
3. **Fixed Scenario:** Scenario 1b-vulnerable configuration only
4. **Observation Quality:** FixedFlatWrapper may lose information

## 6 Conclusion

This study demonstrates that **DQN significantly outperforms PPO and LSTM-PPO** for training Blue team agents in CybORG. The key findings are:

1. Policy gradient methods suffer from **action collapse** - converging to single-action policies regardless of state
2. DQN learns **state-dependent policies** with 39 action changes per episode
3. DQN achieves the best reward (-228.1) with 6 unique actions used deterministically
4. The learned DQN policy implements a **cyclic defense strategy**: Monitor → Analyse → Remove

These results align with the original CybORG paper’s use of DQN-based methods and suggest that value-based RL is more suitable for autonomous cyber defense than policy gradient approaches.

## 7 Future Work

1. **Emulation Validation:** Test trained DQN agent on AWS virtual infrastructure
2. **Multiple Red Agents:** Evaluate against diverse adversary strategies
3. **Observation Engineering:** Investigate alternative state representations

4. **Action Masking:** Only allow actions relevant to compromised hosts
5. **Multi-Agent Training:** Co-evolve Red and Blue agents
6. **Transfer Learning:** Apply trained agents to different scenarios

## References

- [1] Maxwell Standen, Martin Lucas, David Bowman, Toby J. Richer, Junae Kim, and Damian Marriott. *CybORG: A Gym for the Development of Autonomous Cyber Agents*. IJCAI-21 1st International Workshop on Adaptive Cyber Defense, 2021. arXiv:2108.09118.
- [2] John Schulman, Filip Wolski, Prafulla Dhariwal, Alec Radford, and Oleg Klimov. *Proximal Policy Optimization Algorithms*. arXiv preprint arXiv:1707.06347, 2017.
- [3] Volodymyr Mnih, Koray Kavukcuoglu, David Silver, Andrei A. Rusu, Joel Veness, Marc G. Bellemare, Alex Graves, Martin Riedmiller, Andreas K. Fidjeland, Georg Ostrovski, et al. *Human-level control through deep reinforcement learning*. Nature, 518(7540):529–533, 2015.
- [4] Matthew Hausknecht and Peter Stone. *Deep Recurrent Q-Learning for Partially Observable MDPs*. AAAI Fall Symposium Series, 2015.
- [5] Greg Brockman, Vicki Cheung, Ludwig Pettersson, Jonas Schneider, John Schulman, Jie Tang, and Wojciech Zaremba. *OpenAI Gym*. arXiv preprint arXiv:1606.01540, 2016.

## A Hyperparameters

Table 10: Complete Hyperparameter Configuration

Parameter	PPO	LSTM-PPO	DQN
Learning Rate	$3 \times 10^{-4}$	$3 \times 10^{-4}$	$1 \times 10^{-4}$
Batch Size	64	64	64
Discount ( $\gamma$ )	0.99	0.99	0.99
Network	[256, 256]	LSTM(128) + [128, 64]	[256, 256, 128]
Entropy Coef.	0.05	0.01	N/A
Exploration	Entropy	Entropy	$\epsilon$ -greedy
Replay Buffer	N/A	N/A	100,000
Training Steps	500,000	500,000	500,000

## B Action Space Mapping

Table 11: Hierarchical Action Encoding

Action Type	Type Index	Action Range
Monitor	0	0–12
Analyse	1	13–25
Remove	2	26–38
Restore	3	39–51